

SOFTWARE SPECIFICATION

AUTONOMOUS AI AGENTS FOR

COLLABORATIVE DECISION-MAKING

Version: 1.0.0 | **Date:** October 26, 2023 | **Status:** Final Draft

1. SYSTEM OVERVIEW

This specification defines the architecture for a multi-agent system (MAS) capable of autonomous operation and collaborative decision-making. The system enables independent AI agents to perceive their environment, communicate via established protocols, form consensus, and execute coordinated actions to solve complex problems.

2. CORE COMPONENTS

- **Agent Engine:** The runtime environment managing agent lifecycle, state, and local execution loop.
- **Decision Framework:** Logic modules enabling agents to evaluate options using utility functions and probabilistic reasoning.
- **Collaboration Layer:** Messaging infrastructure implementing ACL (Agent Communication Language) for negotiation and consensus.
- **Knowledge Base:** Distributed storage for shared world state, ontologies, and historical decision logs.

3. FUNCTIONAL REQUIREMENTS

- **Autonomy:** Agents must initiate actions without direct human intervention based on assigned goals.
- **Scalability:** System must support dynamic registration and deregistration of 100+ concurrent agents.
- **Interoperability:** Agents must exchange structured messages (JSON/Protobuf) adhering to FIPA standards.
- **Conflict Resolution:** Built-in mechanisms for voting or bidding to resolve resource contention.

- **Auditability:** Full traceability of decision paths and inter-agent communication logs.

4. TECHNICAL ARCHITECTURE

The system follows a microservices-based event-driven architecture.

- **Language:** Python 3.10+ (Core Logic), Go (Messaging Layer).
- **Communication:** gRPC for internal agent RPC; RabbitMQ/Kafka for asynchronous event bus.
- **Persistence:** Redis (Hot State), PostgreSQL (Transactional Data), Vector DB (Semantic Memory).
- **Deployment:** Docker containers orchestrated via Kubernetes.

5. DATA MODELS

Note: Presented in pseudo-schema format.

5.1 Agent Entity

```
struct Agent {  
    id: UUID  
    role: Enum(LEADER, WORKER, OBSERVER)  
    status: Enum(IDLE, BUSY, OFFLINE)  
    capabilities: List<String>  
    utility_function: Function  
    local_state: Map<String, Any>  
}
```

5.2 Task Entity

```
struct Task {  
    id: UUID  
    priority: Integer(1-10)  
    constraints: Map<String, Any>  
    deadline: Timestamp  
    required_capabilities: List<String>  
    assigned_agents: List<UUID>  
}
```

5.3 Decision Record

```
struct Decision {  
    id: UUID  
    context_hash: String  
    options_considered: List<Option>  
    selected_option: Option  
    rationale: String  
    consensus_score: Float  
    timestamp: Timestamp  
}
```

6. APIS & INTERFACES

6.1 Agent Control API (gRPC)

- RegisterAgent(AgentConfig) returns (AgentID)
- UpdateGoal(AgentID, GoalDefinition) returns (Ack)
- GetAgentState(AgentID) returns (StateSnapshot)

6.2 Collaboration Interface (Message Bus)

- **Topic:** agent.negotiation.{proposal_id}
 - Payload: Proposal | CounterProposal | Accept | Reject
- **Topic:** system.broadcast
 - Payload: GlobalStateUpdate | EmergencyHalt

7. AGENT CAPABILITIES

- **Perception:** Agents poll environment state or subscribe to event streams to update their local Knowledge Graph.
- **Reasoning (LLM Integrated):** Usage of chain-of-thought prompting via LLM APIs to generate high-level plans.
- **Learning:** Reinforcement Learning (RL) module updates utility weights based on task success/failure feedback.
- **Protocol Adherence:** Strict enforcement of turn-taking and message formatting during negotiation phases.

8. DECISION-MAKING FRAMEWORK

8.1 Individual Decision

$$\begin{array}{llll} \text{Agents} & \text{maximize} & \text{local} & \text{utility:} \\ U(\text{action}) = w_1 * P(\text{success}) + w_2 * \text{resource_cost} + \\ w_3 * \text{time_efficiency} \end{array}$$

8.2 Collaborative Consensus (Voting)

For group decisions, a weighted voting mechanism is employed:

1. **Proposal Phase:** Leader agent broadcasts a plan.
2. **Evaluation Phase:** Worker agents simulate outcome locally.
3. **Voting Phase:** Agents cast votes (Yes/No/Abstain) weighted by their domain expertise score.
4. **Execution Phase:** If `weighted_yes > threshold`, plan executes.
Otherwise, back to Proposal.

9. PERFORMANCE REQUIREMENTS

| Metric | Target | Condition |
|------------------------|-----------------|--|
| Agent Response Latency | < 200ms | 95th percentile, internal logic processing |
| Consensus Time | < 2 seconds | Groups of up to 10 agents |
| Throughput | 10,000 msgs/sec | Message bus capacity |
| Recovery Time | < 5 seconds | Agent restart after crash |

10. SECURITY & COMPLIANCE

- **Identity:** Mutual TLS (mTLS) for all inter-agent communication.
- **Authorization:** Role-Based Access Control (RBAC) enforced at the API Gateway level.
- **Sandboxing:** Agents execute logic in isolated containers with restricted network egress.
- **Data Protection:** All persistent data encrypted at rest (AES-256).
- **Governance:** "Kill switch" functionality to immediately suspend all agent autonomous actions.